

Gabriel Cárcamo

Senior Backend Engineer — Python & AI

Chile · Remote · UTC-3 / UTC-4

carcamo.gabriel@gmail.com · linkedin.com/in/gabriel-carcamo · Portfolio: brownbull.github.io

SUMMARY

More than 10 years as an engineer at Experian (2014–2025). I built data pipelines in Shell and SQL, later also in Python, and from 2020 deployed and hardened credit-risk models for Experian's banking clients (Chase, Citi, Bank of America) on mainframe, AIX/Linux and AWS, moving and validating data for models that processed 300M+ data points. Since late 2025 I have built LLM features on my own projects, and since April 2026 FastAPI and PostgreSQL backends with typed LLM outputs (Gastify, Gustify), plus a hackathon triage agent with deterministic guardrails and Langfuse tracing. I also built a private set of 22 Claude Code agents, in regular use since January 2026.

BACKEND ENGINEERING · side projects, since April 2026

- **APIs & services:** FastAPI (async endpoints, SQLAlchemy 2.0 async, Alembic) and PostgreSQL 16, with REST API design, built for two product rebuilds, Gastify and Gustify (April–July 2026).
- **Testing & CI:** GitHub Actions with path-filtered pipelines; in Gastify, CI fails if the data-isolation tests (Postgres row-level security) against real Postgres 16 were skipped.

AI ENGINEERING — LLMS & AGENTS · since late 2025

- **Gastify** (expense tracker for the Chilean market; side project, rebuild in progress, test deployment): Gemini output typed through PydanticAI, then an arithmetic check ($\text{items} + \text{tax} - \text{discount} = \text{total}$) that saves the receipt with a review flag when the totals do not add up; a plain-code PDF parser for known statement layouts calls AI only as a consented fallback.
- **SRE Triage Agent** (AgentX Hackathon, built solo with Claude Code, April 7–9, 2026; runs locally): FastAPI + async SQLAlchemy + PostgreSQL; deterministic guardrails before any model call; one structured LLM call per triage, with curated notes on the Solidus e-commerce codebase, through Claude Haiku 4.5 (Anthropic SDK, the default) or Gemini 2.5 Flash (LangChain, optional), plus an experimental engine on the Anthropic Managed Agents API (beta) that explores the repo; every file path the model cites is checked to exist in the repo; one trace per triage in a self-hosted Langfuse, via its SDK.
- **Gustify** (cooking app that suggests recipes from your pantry; side project, in development): FastAPI and PostgreSQL backend; a Gemini recipe generator with schema-typed outputs (google-genai with Pydantic response schemas) is built but switched off in the deployed app, which serves a curated catalog.
- **Agent development:** a private set of 22 Claude Code agents (as of September 2026), in regular use since January 2026 for my own paperwork, research and side projects (tax questions, legal correspondence, audit documentation). Each agent is a YAML definition plus a sidecar of method, knowledge and state, running cases through its own workflow; in 15 of the 22, escalation rules, written as prose instructions, tell the agent to stop and hand the decision to me or refer it to a professional. Deliverables are written for my review, and nothing is published or sent without my confirmation. Internal use; not client work.
- **Retrieval / RAG:** file-based context loading across the agent set, over the Markdown files in its sidecars (selective knowledge loading, pointer-based references to a single source of truth, context compression; no embeddings, vector store or reranking); familiar with the RAG spectrum (naive to agentic, Self-RAG, CRAG, Graph RAG) and its trade-offs; learning FAISS for vector search. *No vector-RAG deployment yet; file-based retrieval to date.*
- **AI-assisted development:** Claude Code is my main coding tool (custom hooks, skills, multi-instance orchestration), inside a test-first, review-gated workflow; earlier, GitHub Copilot at Experian (2024) for model implementation; MCP servers and tools inside Claude Code.

- **Gabe Suite** (self-built Claude Code plugin, since February 2026): 28 skills as of September 2026, including a test-first loop (plan, failing test, code, review, commit, push) with hooks that enforce the order, plus a command center that generates a project's documentation (entities, endpoints, tests, coverage, open work) from git, the test run and the plan files.

INDEPENDENT PROJECTS · 2025 – now

Remote, Chile. Side projects listed above: FastAPI and PostgreSQL backends (Gastify, Gustify) since April 2026, with typed LLM outputs in Gastify, and a solo build for the AgentX Hackathon.

EXPERIENCE — EXPERIAN · 2014–2025

Senior Software Development Engineer

2020 – 2025

Experian · Remote, Chile

- Deployed and hardened credit-risk models for Experian's banking clients (Chase, Citi, Bank of America) on mainframe, AIX/Linux and AWS. The focus was edge cases, test coverage and regulatory requirements, not tuning the models.
- Moved and validated data across platforms for models that processed 300M+ data points, with Python and Shell tools that compare model outputs between mainframe, distributed and cloud runs.
- Cut a model aggregation pipeline from 4+ hours to 15 minutes by restructuring the input for vectorized pandas/numpy group operations and moving to write-once, load-once feature loading.
- Fixed production issues for Experian's banking clients within 24–48 hours. Coordinated 4+ engineering teams on cross-platform implementations, and trained Experian's client teams and internal engineers on migration tooling and coding standards.

Software Developer I & II

2014 – 2020

Experian · Remote, Chile

- Designed and built ETL pipelines in Shell and SQL, adding Python toward the end of this period.
- Administered and optimized DB2, MySQL and PostgreSQL databases.
- Supported critical systems 24/7 for internal and external teams, and built internal tools: Python CLI utilities, Shell libraries and file-format converters.

SKILLS

Languages: Python (about 4 years professional), SQL (11y — DB2, PostgreSQL, MySQL), Shell/Bash (11y), C/C++, JavaScript/TypeScript

Backend (side projects, 2026): FastAPI, SQLAlchemy 2.0 (async), Alembic, PostgreSQL 16, RESTful API design; earlier Flask apps

Platforms (Experian): IBM mainframe (JCL, DB2), AIX/Linux, AWS

AI / LLM: Claude API, Gemini API, PydanticAI, AI agents, prompt engineering, MCP, Langfuse tracing (self-hosted, SDK), OpenTelemetry (console spans), context engineering; RAG (conceptual; learning FAISS)

Cloud / DevOps: AWS (Lambda, EC2, S3; cloud migration work at Experian), Docker, GitHub Actions CI/CD, Git

Data: ETL pipeline design (10+ years), pandas / numpy, PostgreSQL optimization

EDUCATION & LANGUAGES

Ingeniero en Informática (equivalent to a B.S. in Computer Science), Universidad Mayor, Chile, 2016–2020

Spanish: native · **English:** professional fluency (daily working language)